

Team2

1. We have taken Crimedata dataset where 'medv' column is the target column depends on other features in the data. In the first instant We have used SGD optimizer, learning rate = 0.1, batch size = 64 and sigmoid activation function. In the second instant We have used Adam optimizer, learning rate = 0.2, batch size = 32 and tanh activation function. The loss has increased. It may be one of the factor that learning rate has increased so that model tries to train fast and accuracy decreases. The other factors are change in activation function, batch size and optimizer.

Please find the screenshots below



```
from keras.optimizers import Adam
from keras.optimizers import SGD
from sklearn.model_selection import train_test_split
tbc=TensorBoardColab()

df = pd.read_csv('crimedata.csv')
data = pd.DataFrame(df, columns=["crim", "zn", "indus", "chas", "nox", "rm", "age", "dis", "rad", "tax", "ptratio", "b", "lstat", "medv"])
label_col = 'medv'
print(data.describe())

x=data.iloc[:,0:13]
y=data.iloc[:,13]

x_train, x_valid, y_train, y_valid = train_test_split(x, y, test_size=0.3, random_state=87)
np.random.seed(155)

def model1(x_size, y_size):
    model = Sequential()
    model.add(Dense(100, activation="sigmoid", input_shape=(x_size,)))
    model.add(Dropout(0.1))
    model.add(Dense(50, activation="relu"))
    model.add(Dense(20, activation="sigmoid"))
    model.add(Dense(y_size))
    print(model.summary())
    keras.optimizers.SGD(lr=0.1)
    model.compile(loss='mean_squared_error', optimizer=SGD(), metrics=[metrics.mae])
    return(model)

Model = model1(x_train.shape[1], 1)

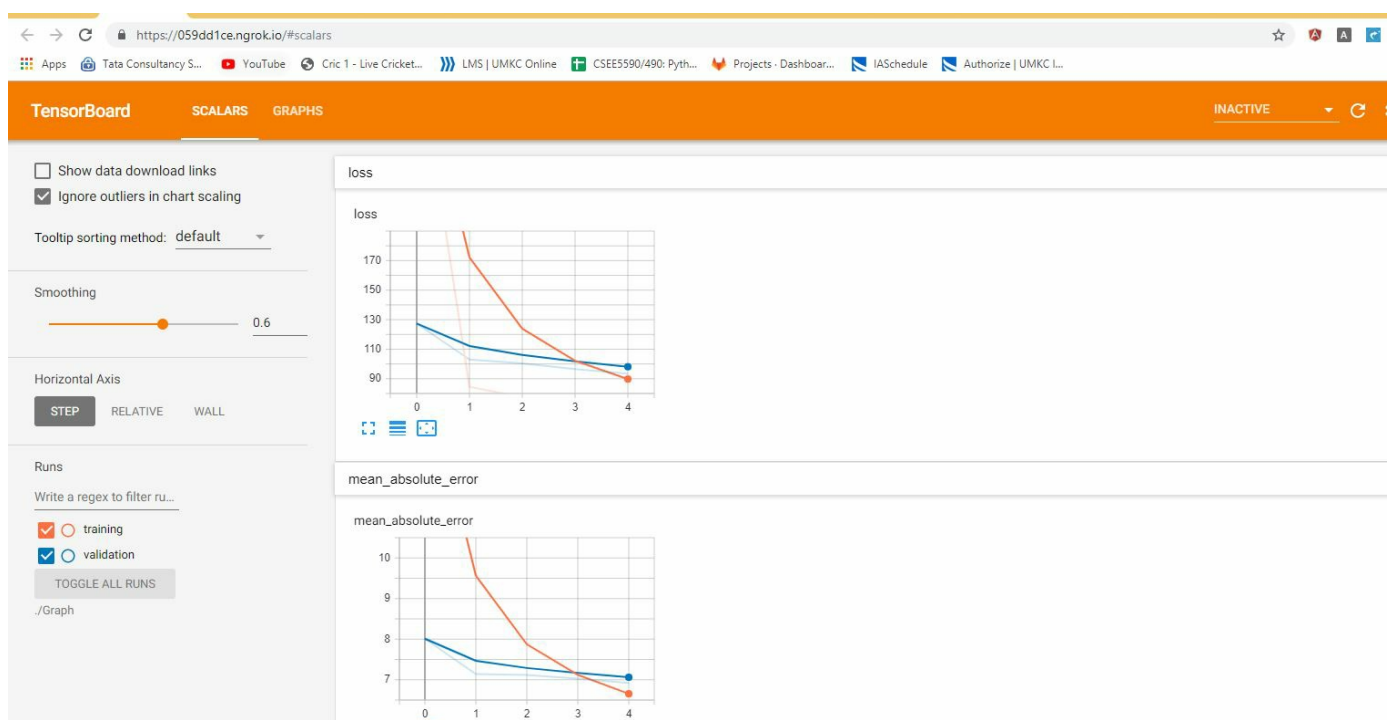

Model.summary()

hist = Model.fit(x_train, y_train, batch_size=64, epochs=5, shuffle=True, verbose=0, validation_data=(x_valid, y_valid), callbacks=[TensorBoardColabCallback(tbc)])

train_score = Model.evaluate(x_train, y_train, verbose=0)
valid_score = Model.evaluate(x_valid, y_valid, verbose=0)

print('Train MAE: ', round(train_score[1], 4), ', Train Loss: ', round(train_score[0], 4))
print('Val MAE: ', round(valid_score[1], 4), ', Val Loss: ', round(valid_score[0], 4))

# accuracy history
plt.plot(hist.history['mean_absolute_error'])
```





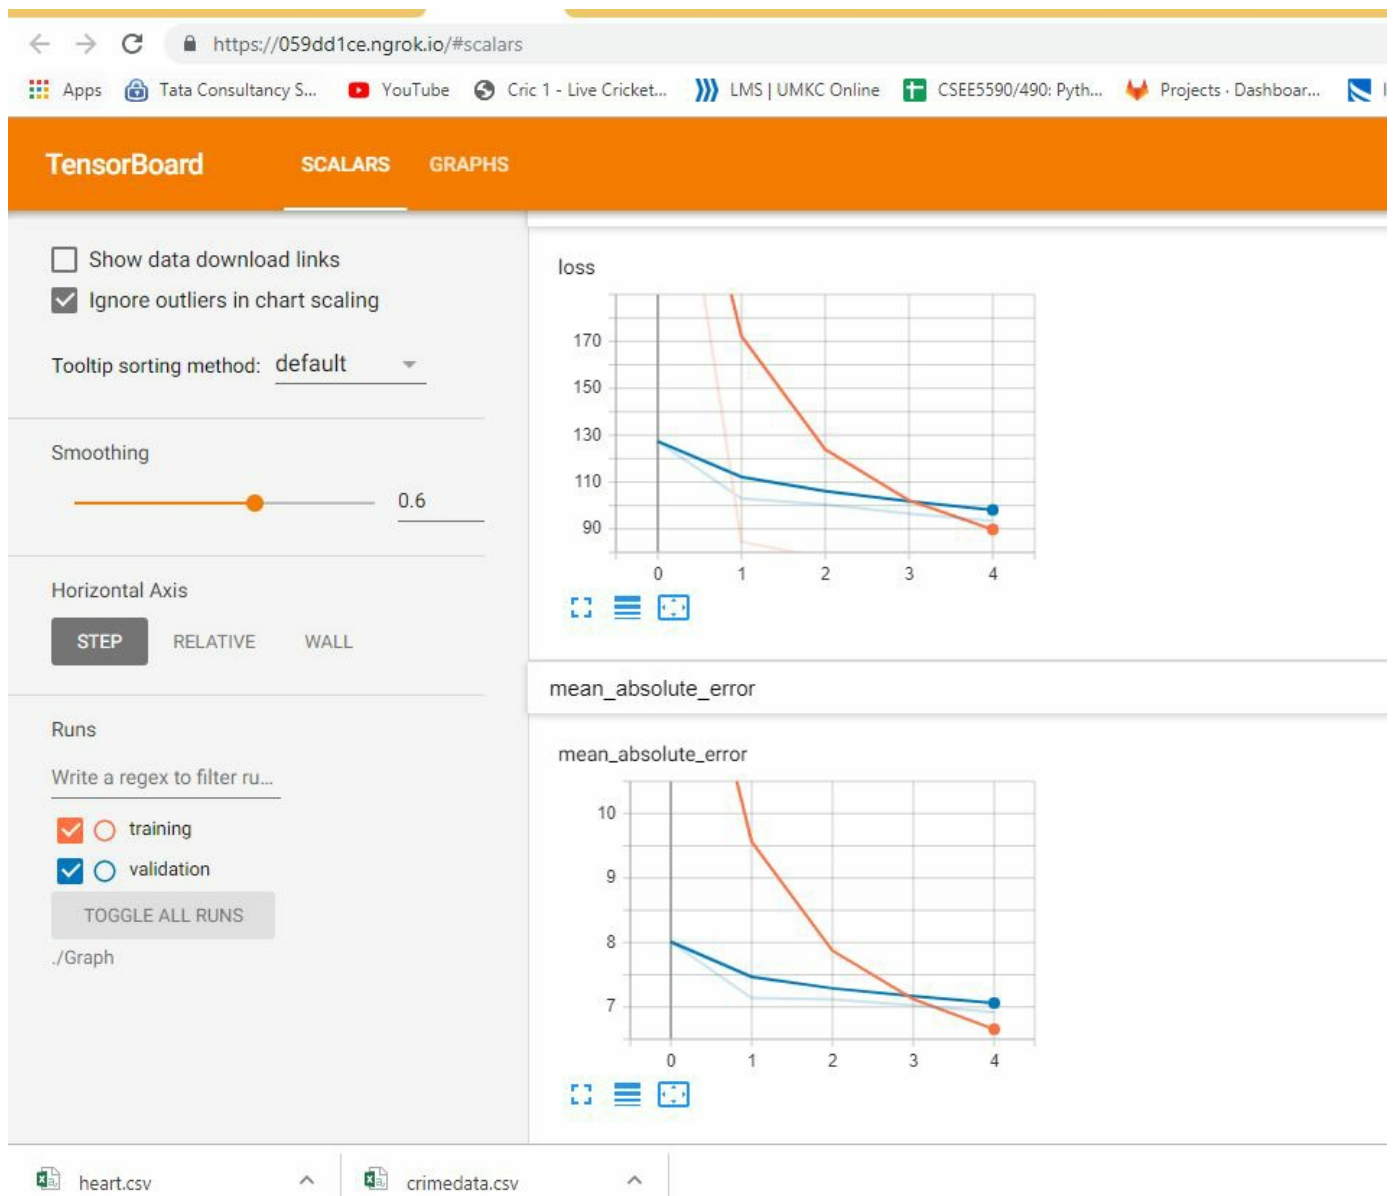
after changing parameters



Train MAE: 15.4861 , Train Loss: 316.5806

Val MAE: 16.4257 , Val Loss: 369.7956





2. We have taken Heart where 'target' column is the target column depends on other features in the data. In the first instant We have used SGD optimizer, learning rate = 0.1, batch size = 256 and sigmoid activation function. In the second instant We have used Adam optimizer, learning rate = 0.2, batch size = 128 and tanh activation function. The accuracy has decreased. It may be one of the factor that learning rate has increased so that model tries to train fast and accuracy decreases. The other factors are change in activation function, batch size and optimizer.

Please find the screenshots below

```
df = pd.read_csv("heart.csv", sep=',')
df.astype(float)
# Normalize values to range [0:1]
df /= df.max()
# split data into features & target columns
x= df.drop(columns = 'target')
y = df['target']
x_train, x_test, y_train, y_test = train_test_split(x, y, train_size = 0.75, test_size = 0.25)
y_train = np_utils.to_categorical(y_train, 2)
y_test = np_utils.to_categorical(y_test, 2)
# Creating the model
model = Sequential()
model.add(Dense(1024, input_shape=(13,), kernel_regularizer = regularizers.l2( 0.01)))
model.add(Activation('relu'))
model.add(Dropout(0.5))
model.add(Dense(1024))
model.add(Activation('tanh'))
model.add(Dropout(0.5))
model.add(Dense(1024))
model.add(Activation('relu'))
model.add(Dropout(0.5))
model.add(Dense(2))
model.add(Activation('sigmoid'))
model.summary()
keras.optimizers.SGD(lr=0.1)

# compile the model
model.compile(loss='categorical_crossentropy', optimizer=SGD(), metrics=['accuracy', keras_metrics.precision(), keras_metrics.recall()])

# train the model

history = model.fit(x_train, y_train, batch_size=256, epochs=10, verbose=0, validation_split=0.25, callbacks=[TensorBoardColabCallback(tbc)])

# make prediction
y_pred = model.predict_classes(x_test)

score = model.evaluate(x_test, y_test, verbose=0)

print('Loss:', score[0])
print('Accuracy:', score[1])
print('Precision:', score[2])
print('Recall:', score[3])
```

dense_2 (Dense)	(None, 1024)	1049600
activation_2 (Activation)	(None, 1024)	0
dropout_2 (Dropout)	(None, 1024)	0
dense_3 (Dense)	(None, 1024)	1049600
activation_3 (Activation)	(None, 1024)	0
dropout_3 (Dropout)	(None, 1024)	0
dense_4 (Dense)	(None, 2)	2050
activation_4 (Activation)	(None, 2)	0
=====		
Total params: 2,115,586		
Trainable params: 2,115,586		
Non-trainable params: 0		

W0723 02:32:07.593160 139751452473216 deprecation.py:323] From /usr/local/lib/python3.6/dist-pack
Instructions for updating:

Use tf.where in 2.0, which has the same broadcast rule as np.where

W0723 02:32:07.824178 139751452473216 deprecation_wrapper.py:119] From /usr/local/lib/python3.6/d

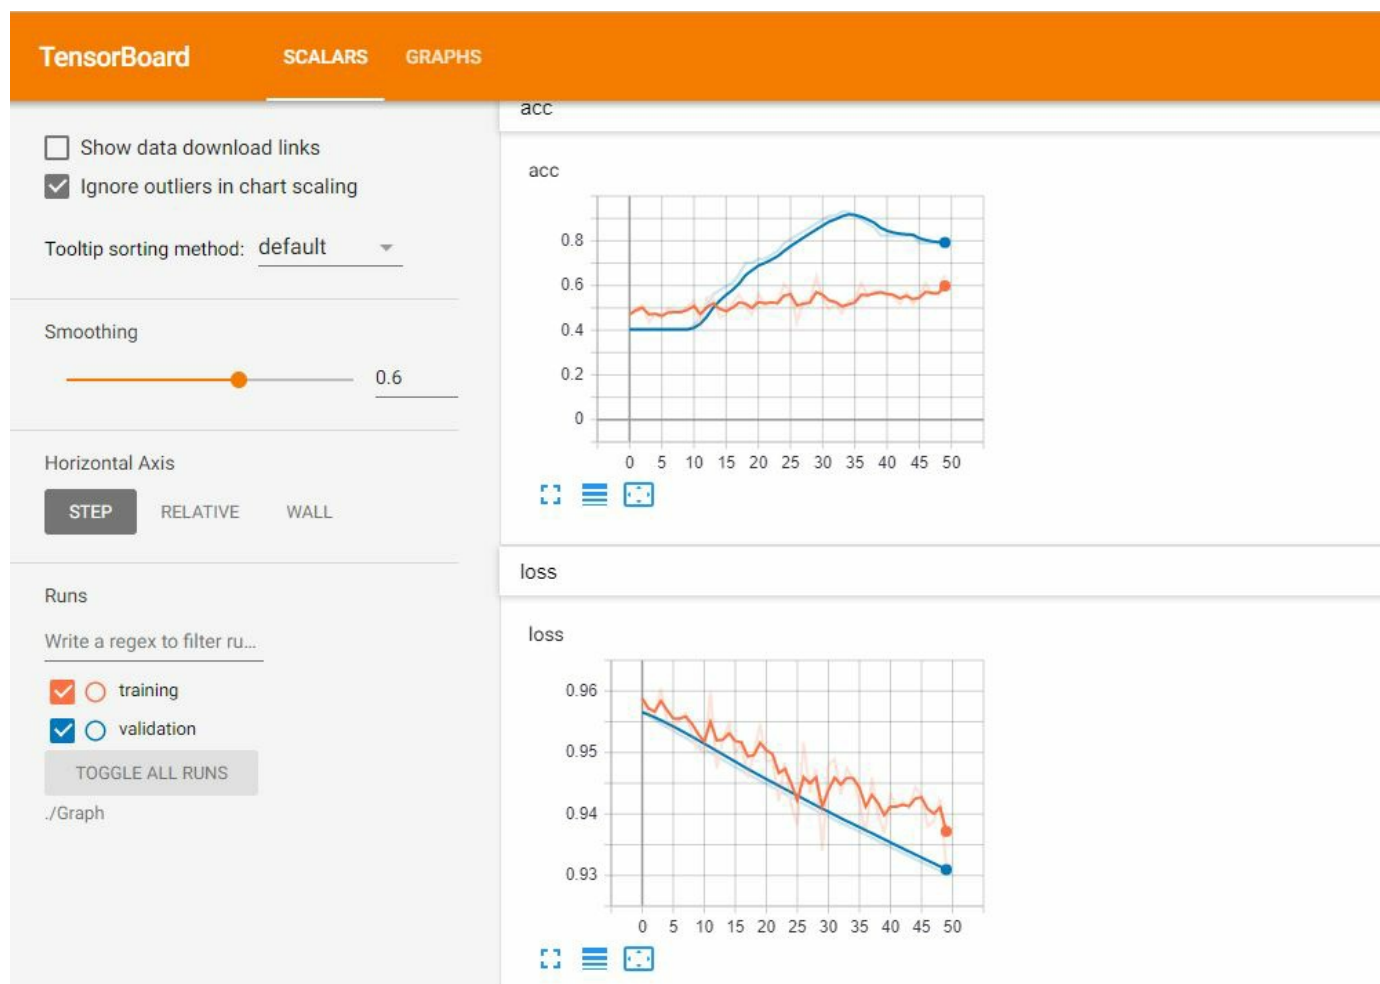
W0723 02:32:07.989934 139751452473216 deprecation_wrapper.py:119] From /usr/local/lib/python3.6/d

Loss: 0.9361106910203633

Accuracy: 0.7368421084002444

Precision: 0.5666666657222222

Recall: 0.9189189164353543



after changing parameters

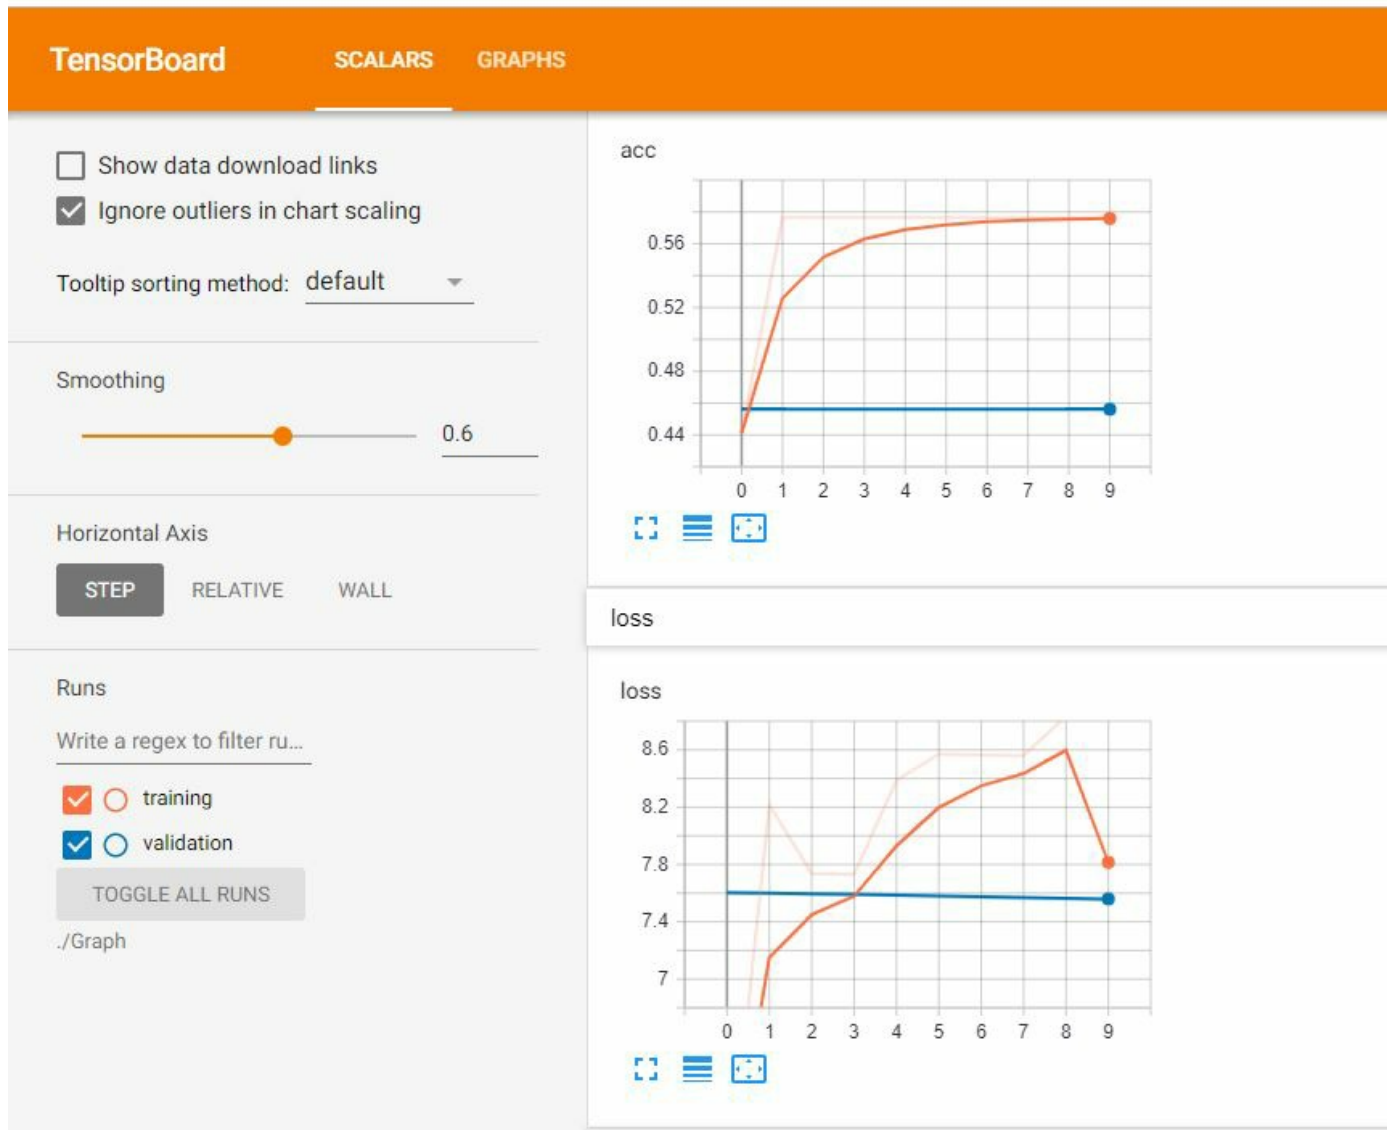
```

▶
📄
dropout_4 (Dropout)      (None, 1024)      0
dense_6 (Dense)          (None, 1024)      1049600
activation_6 (Activation) (None, 1024)      0
dropout_5 (Dropout)      (None, 1024)      0
dense_7 (Dense)          (None, 1024)      1049600
activation_7 (Activation) (None, 1024)      0
dropout_6 (Dropout)      (None, 1024)      0
dense_8 (Dense)          (None, 2)         2050
activation_8 (Activation) (None, 2)         0
=====
Total params: 2,115,586
Trainable params: 2,115,586
Non-trainable params: 0

W0723 02:36:53.066102 140304974346112 deprecation.py:323] From /usr/local/lib/python3.6/d:
Instructions for updating:
Use tf.where in 2.0, which has the same broadcast rule as np.where
W0723 02:36:53.461867 140304974346112 deprecation_wrapper.py:119] From /usr/local/lib/pytl
W0723 02:36:53.753079 140304974346112 deprecation_wrapper.py:119] From /usr/local/lib/pytl

Loss: 8.895012453982705
Accuracy: 0.539473682641983
Precision: 0.46052631639542935
Recall: -0.9999999971428574

```

3. We have taken Spam where 'category' column is the target column depends on text feature in the data using CNN Creating Embedding layer as the input layer with 2000 feature set, input length as the max length of a sentence and 20% of dropouts. Using relu activation function in the Convolution layer of single dimension with 128 neurons with 3 one dimension filter. Adding MaxPooling layer with 2 pool size. Added one more relu activation function Convolution layer with 128 neurons with a feature detector of 2. Adding MaxPooling layer with 2 pool size. Flattening to the single vector. Adding the output layer with sigmoid activation function as there are only 2 categories

(binary Classification).

After building the model - running the model with 5 epochs, 32 batch size.

Checking the train data accuracy, score and the validation data accuracy, score

```
DATA_FILE = 'spam_data.csv'
df = pd.read_csv(DATA_FILE,encoding='latin-1')
print(df.head())
tags = df.Category
texts = df.Message
num_max = 2000
# preprocess
le = LabelEncoder()
cat = le.fit_transform(df.Category)
tok = Tokenizer(num_words=num_max)
tok.fit_on_texts(df.Message)
mat_texts = tok.texts_to_matrix(texts,mode='count')
print(cat[:5])
print(mat_texts[:5])
print(tags.shape,mat_texts.shape)
max_len = 100
cnn_texts_seq = tok.texts_to_sequences(texts)
print(cnn_texts_seq[0])
cnn_texts_mat = sequence.pad_sequences(cnn_texts_seq,maxlen=max_len)
print(cnn_texts_mat[0])
print(cnn_texts_mat.shape)
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(cnn_texts_mat, cat, random_state=42, test_size=.1)
model = Sequential()
model.add(Embedding(2000,20,input_length=max_len))
model.add(Dropout(0.2))
model.add(Conv1D(128, 3, padding='same', activation='relu', kernel_constraint=maxnorm(3)))
model.add(GlobalMaxPooling1D())
model.add(Dense(128))
model.add(Dropout(0.2))
model.add(Activation('relu'))
# model.add(Dense(1))
# model.add(Activation('sigmoid'))
model.add(Dense(1, activation='sigmoid'))
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])

model.fit(X_train, y_train, epochs = 5, batch_size=32, verbose = 1)
# model.fit(X_train, Y_train, epochs = 3, batch_size=batch_size, verbose = 2, callbacks=[TensorBoardColabCallback(tbc)])
score,acc = model.evaluate(X_test,y_test,verbose=1,batch_size=32)
print(score)
print(acc)
print(model.metrics_names)
```

```
0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 49
472 841 756 659 65 8 1328 87 123 352 1329 148 1330 67
58 144]
(5572, 100)
W0723 02:38:53.945644 139960945022848 deprecation_wrapper.py:119] From /usr/local/lib/python3.6/dist-|
W0723 02:38:53.955410 139960945022848 deprecation.py:506] From /usr/local/lib/python3.6/dist-packages,
Instructions for updating:
Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob`.
W0723 02:38:54.038531 139960945022848 deprecation_wrapper.py:119] From /usr/local/lib/python3.6/dist-|
W0723 02:38:54.061208 139960945022848 deprecation_wrapper.py:119] From /usr/local/lib/python3.6/dist-|
W0723 02:38:54.067523 139960945022848 deprecation.py:323] From /usr/local/lib/python3.6/dist-packages,
Instructions for updating:
Use tf.where in 2.0, which has the same broadcast rule as np.where
Epoch 1/5
5014/5014 [=====] - 3s 644us/step - loss: 0.3859 - acc: 0.8624
Epoch 2/5
5014/5014 [=====] - 3s 499us/step - loss: 0.1627 - acc: 0.9356
Epoch 3/5
5014/5014 [=====] - 3s 589us/step - loss: 0.0414 - acc: 0.9886
Epoch 4/5
5014/5014 [=====] - 3s 656us/step - loss: 0.0203 - acc: 0.9942
Epoch 5/5
5014/5014 [=====] - 3s 626us/step - loss: 0.0116 - acc: 0.9972
558/558 [=====] - 0s 291us/step
0.06048352692583342
0.9874551971326165
['loss', 'acc']
```

4. We have done analysis on the same dataset using LSTM model. Created a model and built with LSTM layer. Checking the train data accuracy, score and the validation data accuracy, score

Please find the screenshots below

```

data = pd.read_csv('Spam_Data.csv',encoding='latin-1')
# Keeping only the necessary columns
data = data[['Message','Category']]

data['Message'] = data['Message'].apply(lambda x: x.lower())
data['Message'] = data['Message'].apply((lambda x: re.sub('[^a-zA-z0-9\s]', '', x)))

for idx, row in data.iterrows():
    row[0] = row[0].replace('rt', ' ')

max_fatures = 2000
tokenizer = Tokenizer(num_words=max_fatures, split=' ')
tokenizer.fit_on_texts(data['Message'].values)
X = tokenizer.texts_to_sequences(data['Message'].values)

X = pad_sequences(X)

embed_dim = 128
lstm_out = 196
def createmodel():
    model = Sequential()
    model.add(Embedding(max_fatures, embed_dim,input_length = X.shape[1]))
    model.add(LSTM(lstm_out, dropout=0.2, recurrent_dropout=0.2))
    model.add(Dense(2,activation='softmax'))
    model.compile(loss = 'categorical_crossentropy', optimizer='adam',metrics = ['accuracy'])
    return model
# print(model.summary())

labelencoder = LabelEncoder()
integer_encoded = labelencoder.fit_transform(data['Category'])
y = to_categorical(integer_encoded)
X_train, X_test, Y_train, Y_test = train_test_split(X,y, test_size = 0.33, random_state = 42)

batch_size = 32
model = createmodel()
model.fit(X_train, Y_train, epochs = 3, batch_size=batch_size, verbose = 2)
# model.fit(X_train, Y_train, epochs = 3, batch_size=batch_size, verbose = 2, callbacks=[TensorBo
score,acc = model.evaluate(X_test,Y_test,verbose=2,batch_size=batch_size)
print(score)
print(acc)
print(model.metrics_names)

```



```
print(model.metrics_names)
```

```
Using TensorFlow backend.  
WARNING: Logging before flag parsing goes to stderr.  
W0723 02:39:07.936384 140620413458304 deprecation_wrapper.py:119] From /usr/local/lib/python  
W0723 02:39:07.982941 140620413458304 deprecation_wrapper.py:119] From /usr/local/lib/python  
W0723 02:39:07.989587 140620413458304 deprecation_wrapper.py:119] From /usr/local/lib/python  
W0723 02:39:08.414351 140620413458304 deprecation_wrapper.py:119] From /usr/local/lib/python  
W0723 02:39:08.445332 140620413458304 deprecation.py:506] From /usr/local/lib/python3.6/di  
Instructions for updating:  
Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob`.  
W0723 02:39:09.069054 140620413458304 deprecation_wrapper.py:119] From /usr/local/lib/python  
W0723 02:39:09.145365 140620413458304 deprecation_wrapper.py:119] From /usr/local/lib/python  
W0723 02:39:09.426380 140620413458304 deprecation.py:323] From /usr/local/lib/python3.6/di  
Instructions for updating:  
Use tf.where in 2.0, which has the same broadcast rule as np.where  
Epoch 1/3  
- 45s - loss: 0.1896 - acc: 0.9346  
Epoch 2/3  
- 44s - loss: 0.0492 - acc: 0.9847  
Epoch 3/3  
- 44s - loss: 0.0255 - acc: 0.9925  
0.04851065507759262  
0.9891245241979336  
['loss', 'acc']
```

5. CNN shows more better results compared to LSTM model in case of text classification with large set of data where as in case of small data set, the accuracies and the loss functions are almost same.
The accuracies are more than 98% in case of spam dataset.
6. We have taken monkey species dataset. I am unable to run it in Google collab as dataset is huge to upload to drive. Dataset contains images in the training folder and validation folder. After removing the stopwords tokenized and stemmed with Snowball-Stemmer. Both the datasets after preprocessing are stored in the dictionary. words are converted into the vectors after tokenizing.

CNN model is created and used softmax as the activation function for the last layer. optimizer: Adam() batch size: 256 epochs: 5 Model is evaluated for accuracy.

Please find the screenshots below

```

7 from keras.layers import Dense, Dropout
8 from keras.optimizers import Adam
9 from keras.optimizers import SGD
10 from sklearn.model_selection import train_test_split
11
12 df = pd.read_csv('crimedata.csv')
13 data = pd.DataFrame(df, columns=["crim", "zn", "indus", "chas", "nox", "rm", "age", "dis", "rad", "tax"])
14 label_col = 'medv'
15 print(data.describe())
16
17
18 x_train, x_valid, y_train, y_valid = train_test_split(data.iloc[:,0:13], data.iloc[:,13], test_size=0.2, random_state=155)
19 np.random.seed(155)
20
21 def model1(x_size, y_size):
22     model = Sequential()
23     model.add(Dense(100, activation="tanh", input_shape=(x_size,)))
24     model.add(Dropout(0.1))
25     model.add(Dense(50, activation="relu"))
26     model.add(Dense(20, activation="relu"))
27     model.add(Dense(y_size))
28     print(model.summary())
29     keras.optimizers.SGD(lr=0.1)
30     model.compile(loss='mean_squared_error', optimizer=SGD(), metrics=[metrics.mae])
31     return(model)
32
33 Model = model1(x_train.shape[1], 1)
34
35 Model.summary()
36
37 history = Model.fit(x_train, y_train, batch_size=64, epochs=25, shuffle=True, verbose=0, validation_data=(x_valid, y_valid))
38
39 train_score = Model.evaluate(x_train, y_train, verbose=0)
40 valid_score = Model.evaluate(x_valid, y_valid, verbose=0)
41
42 print('Train MAE: ', round(train_score[1], 4), ' Train Loss: ', round(train_score[0], 4))
43 print('Val MAE: ', round(valid_score[1], 4), ' Val Loss: ', round(valid_score[0], 4))

```

Name	Type	Size	Value
X	int32	(5572, 152)	[[0 0 0 0 ... 68 58 137] [0 0 0 0 ... 444 6 18 ...]
X_len	list	5572	[17, 6, 23, 10, 13, 30, 14, 23, 23, 28, ...]
X_test	int32	(1839, 152)	[[0 0 0 0 ... 249 1788 6] [0 0 0 0 ... 419 12 9 ...]
y_train	int32	(3733, 152)	[[0 0 0 0 ... 282 25 904]

Name	Type	Size	Value
block5_pool (MaxPooling2D)	(None, 7, 7, 512)		0
flatten (Flatten)	(None, 25088)		0
fc1 (Dense)	(None, 4096)		102764544
fc2 (Dense)	(None, 4096)		16781312
drop2 (Dropout)	(None, 4096)		0
fc3 (Dense)	(None, 10)		40970

Total params: 134,301,514
 Trainable params: 40,970
 Non-trainable params: 134,260,544

Epoch 1/100
 137/137 [=====] - 696s 5s/step - loss: 1.3760 - acc: 0.7071 - val_loss: 12.0729 - val_acc: 0.9449
 Epoch 2/100
 137/137 [=====] - 621s 5s/step - loss: 0.3958 - acc: 0.9005 - val_loss: 13.7234 - val_acc: 0.9779
 Epoch 3/100

```

114
115 #####
116
117 # Creating dataframe for validation data in a similar fashion
118 valid_df = []
119 for folder in os.listdir(validation_data):
120     imgs_path = validation_data / folder
121     imgs = sorted(imgs_path.glob('*.jpg'))
122     for img_name in imgs:
123         valid_df.append((str(img_name), labels_dict[folder]))
124
125 valid_df = pd.DataFrame(valid_df, columns=['image', 'label'], index=None)
126 # shuffle the dataset
127 valid_df = valid_df.sample(frac=1.).reset_index(drop=True)
128
129 #####
130
131 # How many samples do we have in our training and validation data?
132 print("Number of training samples: ", len(train_df))
133 print("Number of validation samples: ", len(valid_df))
134
135 # sneak peek of the training and validation dataframes
136 print("\n", train_df.head(), "\n")
137 print("=====\n")
138 print("\n", valid_df.head())
139
140
141
142
143 # some constants(not truly though!)
144
145 # dimensions to consider for the images
146 img_rows, img_cols, img_channels = 224,224,3
147
148 # batch size for training
149 batch_size=8
150

```

Name	Size	Type	Created
Screenshots	File Folder		22-07-2019 17:56
model1	512.6 MB File		22-07-2019 19:22

Name	Type	Size	Value
block5_conv2 (Conv2D)	(None, 14, 14, 512)		2359808
block5_conv3 (Conv2D)	(None, 14, 14, 512)		2359808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)		0
flatten (Flatten)	(None, 25088)		0
fc1 (Dense)	(None, 4096)		102764544
fc2 (Dense)	(None, 4096)		16781312
drop2 (Dropout)	(None, 4096)		0
fc3 (Dense)	(None, 10)		40970

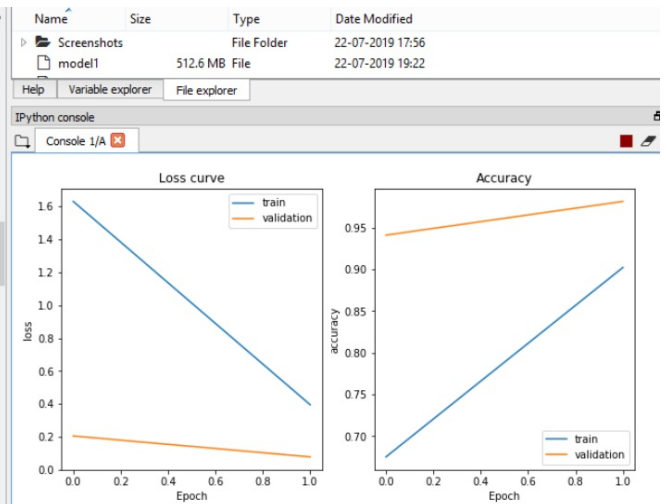
Total params: 134,301,514
 Trainable params: 40,970
 Non-trainable params: 134,260,544

Epoch 1/2
 137/137 [=====] - 743s 5s/step - loss: 1.6308 - acc: 0.6752 - val_loss: 0.2044 - val_acc: 0.9412
 Epoch 2/2
 137/137 [=====] - 628s 5s/step - loss: 0.3945 - acc: 0.9024 - val_loss: 0.0774 - val_acc: 0.9816

```

114
115 #####
116
117 # Creating dataframe for validation data in a similar fashion
118 valid_df = []
119 for folder in os.listdir(validation_data):
120     imgs_path = validation_data / folder
121     imgs = sorted(imgs_path.glob('*.jpg'))
122     for img_name in imgs:
123         valid_df.append((str(img_name), labels_dict[folder]))
124
125 valid_df = pd.DataFrame(valid_df, columns=['image', 'label'], index=None)
126 # shuffle the dataset
127 valid_df = valid_df.sample(frac=1.).reset_index(drop=True)
128
129 #####
130
131 # How many samples do we have in our training and validation data?
132 print("Number of training samples: ", len(train_df))
133 print("Number of validation samples: ", len(valid_df))
134
135 # sneak peek of the training and validation dataframes
136 print("\n", train_df.head(), "\n")
137 print("=====\n")
138 print("\n", valid_df.head())
139
140
141
142
143 # some constants(not truly though!)
144
145 # dimensions to consider for the images
146 img_rows, img_cols, img_channels = 224,224,3
147
148 # batch size for training
149 batch_size=8
150

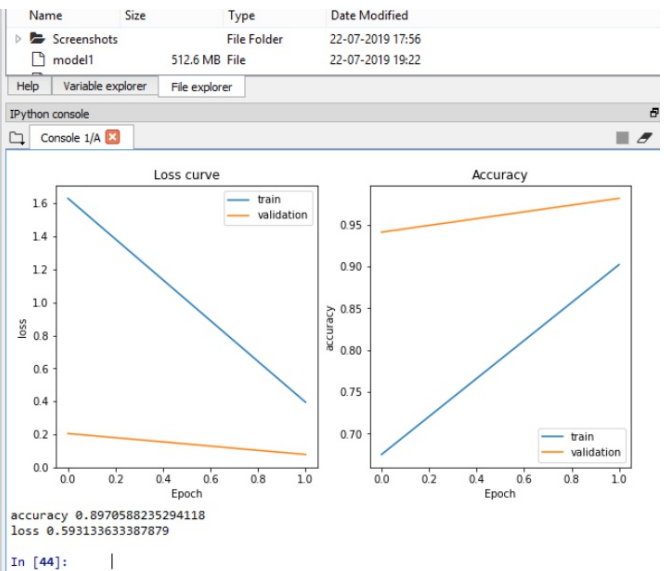
```



```

114
115 #####
116
117 # Creating dataframe for validation data in a similar fashion
118 valid_df = []
119 for folder in os.listdir(validation_data):
120     imgs_path = validation_data / folder
121     imgs = sorted(imgs_path.glob('*.jpg'))
122     for img_name in imgs:
123         valid_df.append((str(img_name), labels_dict[folder]))
124
125 valid_df = pd.DataFrame(valid_df, columns=['image', 'label'], index=None)
126 # shuffle the dataset
127 valid_df = valid_df.sample(frac=1.).reset_index(drop=True)
128
129 #####
130
131 # How many samples do we have in our training and validation data?
132 print("Number of training samples: ", len(train_df))
133 print("Number of validation samples: ", len(valid_df))
134
135 # sneak peek of the training and validation dataframes
136 print("\n", train_df.head(), "\n")
137 print("=====\n")
138 print("\n", valid_df.head())
139
140
141
142
143 # some constants(not truly though!)
144
145 # dimensions to consider for the images
146 img_rows, img_cols, img_channels = 224,224,3
147
148 # batch size for training
149 batch_size=8
150

```



Thank you